

Project Proposal: Machine Learning System for Clinical Diabetes Risk Prediction

Project Guide: Pushpen Lahiri

Project Team:

Name	Examination Roll No.
Ranadeep Saha	11600224078
Rettika Hati	11600224079
Sanchita Banerjee	11600224088
Sathi Basak	11600224091
Shreetama Senapati	11600224105
Shreya Das	11600224106

1. Executive Summary

Diabetes mellitus is a growing global health crisis requiring early diagnostic intervention. This project proposes the development of an end-to-end machine learning pipeline capable of predicting a patient's diabetes risk using a modern clinical and demographic dataset. Utilizing the 100,000-row *Diabetes Prediction Dataset*, this system will serve as an automated screening tool. The project will evaluate a linear statistical baseline (**Logistic Regression**) against an optimized ensemble model (**Random Forest Classifier**). Rather than optimizing blindly for overall accuracy, this implementation introduces a customized misclassification cost matrix to successfully establish clinical equilibrium between diagnostic Precision and Recall.

2. Project Objectives

- **Heterogeneous Data Engineering:** Build an efficient preprocessing pipeline using column-specific transformations to elegantly handle a mix of categorical demographic factors and continuous biometric indicators.
- **Algorithmic Class Stabilization:** Systematically mitigate a severe ~91.5% to 8.5% class imbalance by moving away from uniform loss functions and incorporating explicit class-weighting penalties.
- **Comparative Machine Learning Architecture:** Evaluate the mathematical boundaries of a baseline Logistic Regression model against a non-linear Random Forest Classifier.
- **Clinical Metric Optimization:** Establish a highly calibrated decision framework targeting a precision threshold of ≥ 72 and a recall sensitivity of ≥ 75 for positive diabetic detection to minimize hazardous diagnostic false negatives.
- **Interoperable Deployment Layout:** Design the final pipeline architecture for seamless export, preparing the serialized model for a lightweight, interactive web application framework.

3. Methodology & Technical Workflow

3.1 Advanced Pre-processing & Feature Encoding

To process complex patient records without loss of data fidelity, the pipeline utilizes a modular, multi-track feature transformation design:

- **Categorical Feature Pipeline:** Variables such as gender and smoking_history are parsed and converted into structural numeric arrays using One-Hot Encoding to avoid arbitrary ordinal bias.
- **Numerical Feature Pipeline:** Clinical metrics such as age, BMI, HbA1c_level, and blood_glucose_level are standardized using a standard scaling vector to eliminate feature dominance caused by disparate numerical scales:

$$z = (x - \mu) / \sigma$$

- **Unified Pipeline execution:** Both pipelines are integrated into a vectorized ColumnTransformer to guarantee zero data leakage between the 80% training set and 20% validation testing split.

3.2 Imbalance Remediation via Tailored Cost Matrix

Standard classifiers often fall victim to the "Accuracy Paradox" when faced with a 9:1 healthy-to-diabetic class distribution. This proposal completely rejects generic classification frameworks in favor of a **Custom Class-Weight Matrix**.

The final Random Forest engine assigns a strict **1:5 penalty ratio** (class_weight={0: 1, 1: 5}). By mathematically penalizing misclassified minority instances five times more severely than majority instances, the algorithm shifts its internal decision boundaries away from conservative baselines toward a balanced, clinically safe diagnostic territory.

3.3 Algorithmic Optimization Strategy

The model comparison architecture establishes a strict performance escalation path:

- **Baseline Phase (Logistic Regression):** Tracks the log-odds of the binary target. While computationally efficient, linear limits typically constrain recall performance over highly interactive data distributions.
- **Ensemble Phase (Random Forest Classifier):** Utilizes an ensemble of uncorrelated decision trees optimized with a targeted structural depth (max_depth=10) and dynamic tree allocations (n_estimators=50). This prevents unmitigated tree depth variance while capturing complex interactions between chronic biomarkers like HbA1c_level and blood_glucose_level.

4. Rigorous Evaluation Metrics

To validate clinical readiness, model success is evaluated through a strict classification matrix rather than global metric accuracy:

- **Diabetic Recall (Sensitivity):** Aiming for a minimum targets of ≥ 75 , ensuring that the vast

majority of true diabetic individuals are successfully captured.

- **Diabetic Precision:** Aiming to maintain a threshold of ≥ 72 to prevent excessive false alarms from disrupting clinical resources.
- **F1-Score Evaluation:** Maximizing the harmonic mean of precision and recall to verify model stability within the highly imbalanced cohort.
-

5. Implementation Roadmap

Phase	Milestone	Operational Goals
Phase 1	Pipeline Construction	Structured schema parsing, ColumnTransformer building, and 80/20 train-test splitting.
Phase 2	Model Tuning & Matrix Calibration	Implementing baseline Logistic Regression followed by hyperparameter tuning of the Random Forest using a 1:5 loss-weight ratio.
Phase 3	Diagnostic Validation	Generating comprehensive classification reports, extracting macro/weighted F1-scores, and evaluating model bounds.
Phase 4	Deployment Integration	Serializing the preprocessor and optimized model using joblib for immediate integration into an interactive web UI dashboard.

6. Project Impact & Significance

This project demonstrates how rigorous data engineering and purposeful parameter adjustments can transform an highly skewed 100,000-row health dataset into a balanced, medically reliable diagnostic engine. By selecting a highly specialized algorithmic strategy, this system bridges the gap between pure computer science theory and safe, deployable health informatics.